

# Technical Summary of the Program

This project implements a Transformer-based factor prediction pipeline for quantitative equity modeling, integrating data preprocessing, sequence construction, model training, and rolling prediction. The outcome of the model is the **trading signal of a portfolio**.

## 1. Data Processing (dataproc.py)

This module handles raw financial data transformation. It standardizes numeric features, manages missing values, applies outlier clipping, and constructs factor tensors suitable for sequence models. It builds both *cross-sectional* and *temporal* input representations, aligning with the factor definitions used in the research. It also implements rolling-window sampling logic, allowing dynamic dataset updates as new market data arrives.

## 2. Model Architecture (model\_utils.py)

The model utilities define the core Transformer and Attentional LSTM (ALSTM) architectures for multi-horizon return prediction. The Transformer encoder employs multi-head scaled dot-product attention with feedforward layers, layer normalization, and dropout regularization. Custom modules such as *Weighted MSE* and *IC-based loss functions* are included to directly optimize correlation between predicted and realized returns.

## 3. Training Pipeline (train\_pip.py)

This script integrates model initialization, optimizer setup, and full-cycle training with validation and early stopping. It supports *rolling-window training* to simulate realistic out-of-sample forecasting. Model checkpoints and best-performing weights are saved automatically. The pipeline also tracks performance metrics such as RankIC, ICIR, and Sharpe ratio for continuous evaluation.

## 4. Main Execution and Testing (main.py, test.py)

main.py serves as the experiment entry point, controlling all configuration parameters. It orchestrates the rolling retraining procedure across trading months or weeks. test.py validates trained models on hold-out periods and generates predicted factor values for downstream portfolio construction and backtesting.

## Limitations

- The model does not rebuild training sets progressively within each month's loop — therefore, it's not fully aligned with a rolling backtesting framework, making it less suitable for "timing retraining."
- Missing value handling remains debatable.

## Workflow

### Writing new loss function


Loss function

```

Python
squared_error = torch.pow(y_pred - y_true, 2)
weights = torch.exp(y_true)
weighted_loss = weights * squared_error

```

Outcome:



However, extending the backtest still showed that high ranks (top 1–2 groups) sometimes underperformed.

The exponential weighting created excessive loss gaps between G1 and G2.

So, Two versions were tested:

```
Sigmoid(x) = 1 / (1 + torch.exp(-x))
```

```
z = 40 * (x - 0.8)
```

```
y = 1 / (1 + torch.exp(-z))
```

```
y_fixed = 0.1 + 0.9 * y.
```

```
z = 15 * (x - 0.5)
```

```
y = 1 / (1 + torch.exp(-z))
```

```
y_fixed = 0.5 + 0.5 * y.
```

The **second version performed better**, lifting the lower bound while extending the upper plateau. But after that I directly use the IC as the loss function.


### Parameter Finding

Earlier conclusions may be less reliable; see *Transformer Training* for specific backtest outcomes.

- Best lookforward = 10 days
- Best sequence length (seq) = 5 (may change for deeper networks)
- IC Loss significantly outperforms MSE Loss
- For IC Loss, larger batch sizes (1k–20k) yield more stable IC computation
- Small networks may overfit after 30 epochs (val set not fully out-of-sample)
- Longer lookforward horizons perform better

## Optimal Parameters Combination(AMP enabled)

```
Python
model_size = 'ori'
#amp = True
date_start = None
date_end = None
date_start = "2018-01-01"
date_end = "2024-12-31"
n_epochs = 30
lr = 1e-5
patience = 20
valid_days = 0
lookback_grid = 252
seq_len = 5
use_amp = True
batch_size=20000
```



## AMP experiment

While the need for FP32 master weights is not universal, there are two possible reasons why a

number of networks require it. One explanation is that updates (weight gradients multiplied by the learning rate) become too small to be represented in FP16 - any value whose magnitude is smaller than  $-24$  becomes zero in FP16. We can see that approximately 5% of weight gradient values have exponents smaller than  $-24$ . These small valued gradients would become zero in the optimizer when multiplied with the learning rate and adversely affect the model accuracy. Using a single-precision copy for the updates allows us to overcome this problem and recover the accuracy.

- While not all networks require FP32 master weights, many do for two main reasons:
- Updates (weight gradients learning rate) may be too small for FP16 representation.
- About 5% of weight gradients have exponents  $< 24$  become zero in FP16 accuracy loss.
- Maintaining a single-precision copy for updates resolves this issue.

#### FP16 vs FP32 Computations

- Converted to FP16: nn.Linear, matmul
- Kept in FP32 (numerically sensitive):
- softmax, LayerNorm, BatchNorm, reductions (sum, mean), and loss

## Speed Improvement Test (5 Epochs, 24th Month)

```
seq_len=5,
n_epochs=5,
valid_days=0,
train_model_no_val
batch_sizes=(512,1024,2048,4096),
amp_options=(False, True),
repeats=3,
split_index=23,
criterion=ICLoss()
```

Default model sizes:

batch_size	use_amp	repeats	avg_time_sec	std_time_sec	min_time_sec	max_time_sec
512	False	5	37.778353	0.397365	37.191360	38.427222
512	True	5	38.888162	0.310285	38.467932	39.173111
1024	False	5	44.276889	0.434376	43.649449	44.820548
1024	True	5	44.580894	0.701683	43.872169	45.877156
2048	False	5	47.029617	0.199370	46.801178	47.302885
2048	True	5	47.051206	0.352638	46.632982	47.642976
4096	True	5	58.678543	0.474201	57.896710	59.371733
4096	False	5	58.945991	0.315233	58.454034	59.314086

### Medium

```
d_model=128, num_heads=4, ff_dim=512, num_layers=4, seq=5
```

===== Benchmark Summary (split #24) =====						
batch_size	use_amp	repeats	avg_time_sec	std_time_sec	min_time_sec	max_time_sec
1024	False	3	48.363309	0.186900	48.203201	48.625493
1024	True	3	48.545183	0.352808	48.209780	49.032790
512	False	3	52.888417	0.648896	52.290412	53.790238
512	True	3	56.642478	0.190116	56.400383	56.864810
256	False	3	99.624951	2.860424	97.517334	103.668994
256	True	3	110.293901	0.499010	109.658346	110.877328

## Large

d\_model=256, num\_heads=8, ff\_dim=1024, num\_layers=6, seq=5

===== Benchmark Summary (split #24) =====						
TransformerNet(input_dim=len(factor_cols), d_model=256, num_heads=8, num_layers=6, ff_dim=1024)						
batch_size	use_amp	repeats	avg_time_sec	std_time_sec	min_time_sec	max_time_sec
512	False	3	75.769370	0.087357	75.656439	75.869214
512	True	3	73.196284	0.226641	72.894548	73.440776
1024	True	3	57.019442	0.141353	56.872506	57.210292
1024	False	3	68.941256	0.080487	68.875533	69.054601
2048	True	3	68.246054	0.755563	67.277876	69.121663
2048	False	3	68.080577	2.408195	65.850089	71.424678
4096	True	3	86.149766	1.154961	85.111487	87.760870
4096	False	3	87.190958	0.560415	86.656037	87.964868

d\_model=256, num\_heads=8, ff\_dim=1024, num\_layers=6, seq=21

===== Benchmark Summary (split #24) =====						
TransformerNet(input_dim=len(factor_cols), d_model=256, num_heads=8, num_layers=6, ff_dim=1024) seq =21						
batch_size	use_amp	repeats	avg_time_sec	std_time_sec	min_time_sec	max_time_sec
1024	True	1	83.357674	0.0	83.357674	83.357674
1024	False	1	151.712116	0.0	151.712116	151.712116
2048	False	1	139.938540	0.0	139.938540	139.938540
2048	True	1	78.905366	0.0	78.905366	78.905366
4096	False	1	129.795845	0.0	129.795845	129.795845
4096	True	1	82.101910	0.0	82.101910	82.101910

===== Benchmark Summary (split #24) =====						
batch_size	use_amp	repeats	avg_time_sec	std_time_sec	min_time_sec	max_time_sec
1200	False	3	151.402769	0.146969	151.203169	151.552764
1200	True	3	89.216995	0.138210	89.109503	89.412116
1500	False	3	147.265340	1.288445	145.762713	148.909244
1500	True	3	85.905587	1.080236	84.530604	87.169648
1800	False	3	149.424227	0.281072	149.052962	149.732839
1800	True	3	87.453449	0.208405	87.268903	87.744734

## Conclusion

- AMP yields **significant speedup** in large networks.
- In small networks, AMP can **increase** computation time instead.